>>> Introduction to SQL
>>> Featuring MySQL and T-SQL

Daniel Fryer [†]
Feb, 2023

[†]daniel@vfryer.com

Timetable				
9:00am	-	10:30am	lecture 1	(1.5 hr)
10:30am	-	11:00am	morning tea	(30 min)
11:00am	-	12:30pm	lecture 2	(1.5 hr)
12:30pm	-	1:30pm	lunch	(1 hr)
1:30pm	-	3:00pm	guided exercises	(1.5 hr)
3:00pm	-	5:00pm	one-on-one help	(2 hr)

Day 1

Page is hyperlinked: click a topic above to jump to it.

* S. Q. L. (Structured Query Language)
* 'SEQUEL' (Structured English Query Language)

We will be boldly using two dialects of (ISO/ANSI) SQL:

* T-SQL (Microsoft, proprietary)
* MySQL (Oracle, open-source)

>>> No ice breaker... yet

Breakout rooms during the exercises

Past experience

Definition

A Kahoot is a fun group quiz that we'll do here and there throughout the course. Join in to test your skills.

And now for a practice Kahoot…

```
>>> Where are we now?
```

Day 1
1. Introduction
2. Intro to relational model
3. Tables and relationships
4. Programming in SQL
5. Basic SQL
6. Joining in SQL

Page is hyperlinked: click a topic above to jump to it.

>>> Let the learning begin

A Relational Database Management System (RDBMS).

Definition
A DBMS is a large collection of interdependent programs all working together to define, construct, manipulate, protect and otherwise manage a database. An RDBMS is the most popular kind of DBMS.

SQL is a programming language for talking to your RDBMS.

The most popular Relational Database Management Systems
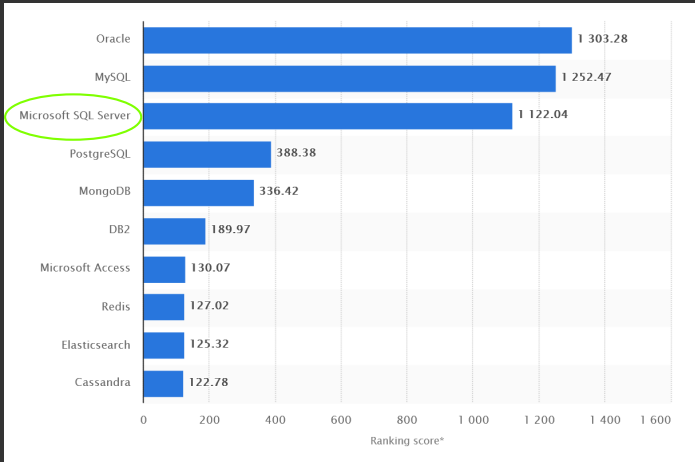
Source: statista.com

Data Analyst

Responsibilities

-Curate the data
-Visualize and report data

Tools
- Excel
- SQL
- Tableau

Skills
- Analytics
- Communication & Visualization

Business Facing
- Sometimes

Salary
$65,000

Data Scientist

Responsibilities

-Source data
-Analyze data
-Run experiments

-Build models
-Recommend solutions
-Storytell

Tools
- Python
- R
- SQL

Skills
- Analytics
- Communication
- Story-tell
- Model building
- Math
- Coding

Business Facing
- Yes

Salary
$120,000

Data Engineer

Responsibilities

-Build data pipelines and warehouse
-Manage scalability of data products

Tools
- Java
- C++
- Kubernetes
- Hadoop
- Spark
- Python

Skills
- Coding
- Model implementation
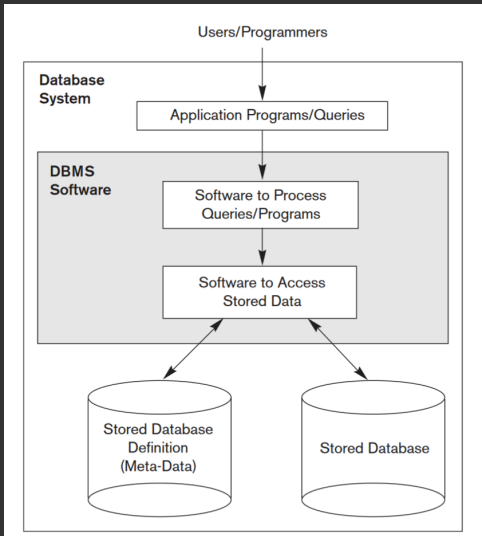
Business Facing
- Not Really

Salary
$110,000

*Matha*Magicians

A layer of abstraction between human and machine

```
>>> RDBMS
```

Grandfather of SQL and RDBMS, in the 1970s:
    ''Future users of databases should be protected
    from having to know how the data is organised in the
    machine.'' - Ted Codd (IBM researcher).

To talk to humans and machines, the RDBMS should have a model of the world that is intuitive to both. This model is called the Relational Model.

Intuition

The Relational Model is the 'common tongue' between the humans and the machines. It has a nice formal mathematical definition, so it is easy for machines to work with. For the humans, it has a simple intuitive description in terms of tables and relationships between tables!

Why does SQL not make it easy to work with CSVs?

As a language for working with structured data stored in relational databases, SQL is not optimized for working with CSV files which are flat, unstructured, and lack a well-defined schema. CSV files do not have the same constraints and relationships that are inherent in a relational database, making it more difficult to work with them in SQL. Additionally, SQL requires data to be in a specific format, while CSV files can vary widely in format, making it challenging to parse and manipulate the data in a consistent manner. SQL also lacks built-in functions for handling CSV data, requiring the use of external tools or custom code. Overall, these factors contribute to making it more difficult to work with CSV files in SQL.

Now can you explain it like one pirate talking to another?

Arrr matey, SQL be designed fer structured data, ya see? But CSVs be flat and lackin' a proper schema, so workin' with 'em in SQL be a challenge. Ye can't just plop 'em in and expect everything to fit nicely like ye would with a proper database. And SQL be strict with its format, but CSVs be all over the map, makin' it tough to parse and manipulate the data. It be lackin' the proper tools too, so ye have to use external tricks or write yer own code. It be a rough sea out there fer workin' with CSVs in SQL.

Okay now explain it like I'm 5, but make it sound like a passage from the King James Bible

Verily, I say unto thee, SQL is a tool for organizing data in a certain manner, even as ye would put thy toys in a chest. But CSVs, they be different, like a pile of sticks thrown upon the ground. They lack a proper order and structure, making it hard for SQL to understand them. SQL desireth a specific way of doing things, but CSVs change and be different every time. Thus, it requireth extra effort and tools to make them work with SQL, even as one must use a hammer to nail a board.

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

>>> What's the takeaway from all this??

When using SQL, you'll always be working with tables. This is
(deceptively) simple and intuitive. Underlying that, there is
a really powerful system that let's you talk to the machine
in a fairly ideal way. This makes SQL very efficient.

The tradeoff? Some parts of SQL will be really simple and
intuitive. Others can at first be frustrating and confusing.
A little practice goes a loooooong way.

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Table name

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Row
(record)

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Column (attribute)

Column names (attribute names)

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Primary key	Friends		
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

Primary key	Friends		
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

* Every table should have a primary key
* No two rows can have the same entry
* There must be no NULL entries

Friends(FriendID, FirstName, LastName, FavColour)

```
>>> One more thing: The data types of attributes
```

Friends (FriendID, FirstName, LastName, FavColour)
              int

Definition
An integer is a positive or negative whole number.

Friends(FriendID, FirstName, LastName, FavColour)
varchar   varchar   varchar

Definition
Varchar stands for 'variable length character.'
It is a string of characters of undetermined length.

```
>>> Where are we now?
```

Day 1
 1. Introduction
 2. Intro to relational model
 3. Tables and relationships
 4. Programming in SQL
 5. Basic SQL
 6. Joining in SQL

Page is hyperlinked: click a topic above to jump to it.

1. One-to-many relationships
2. Primary and foreign keys
3. Many-to-many relationships
4. One-to-one relationships

```
>>> One-to-many relationships
```

* For each car there are *many* wheels.

* For each car there are *many* wheels.

  * For each car there are *many* wheels.
    But each wheel belongs to only *one* car.

* For each car there are *many* wheels.
  But each wheel belongs to only *one* car.
* One bank can have *many* accounts.
  But each account belongs to *one* bank.

* For each friend there are *many* pets.
  But each pet belongs to only *one* friend.

Where do we put the extra pets?

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

* For each friend there are *many* pets.
  But each pet belongs to only *one* friend.

Where do we put the extra pets?

Friends				
FriendID	FirstName	…	PetName$_1$	PetName$_2$
1	X	…	NULL	NULL
2	Y	…	Chikin	NULL
3	Z	…	Cauchy	Gauss

Ideas?

Friends				
FriendID	FirstName	…	$PetName_1$	$PetName_2$
1	X	…	NULL	NULL
2	Y	…	Chikin	NULL
3	Z	…	Cauchy	Gauss

* Have to store NULL in every entry with no pet

 * Have to store NULL in every entry with no pet
 * What if I meet a friend with 3+ pets? Many more NULLs

* Have to store NULL in every entry with no pet
* What if I meet a friend with 3+ pets? Many more NULLs
* New one-to-many relationship between pets and toys?

* Have to store NULL in every entry with no pet
* What if I meet a friend with 3+ pets? Many more NULLs
* New one-to-many relationship between pets and toys?
* Pets are tied to owners. Delete an owner → delete pets

  * Have to store NULL in every entry with no pet
  * What if I meet a friend with 3+ pets? Many more NULLs
  * New one-to-many relationship between pets and toys?
  * Pets are tied to owners. Delete an owner $\rightarrow$ delete pets
  * Ambiguity. Is information related to pets or owners?

So what do we do instead?

Suspense.
Kahoot time! The relational model.

>>> What if we do this instead?

Friends			
FriendID	FirstName	…	PetName
1	X	…	NULL
2	Y	…	Chikin
3	Z	…	Cauchy
3	Z	…	Gauss

```
>>> What if we do this instead?
```

Friends			
FriendID	FirstName	…	PetName
1	X	…	NULL
2	Y	…	Chikin
3	Z	…	Cauchy
3	Z	…	Gauss

This causes data redundancy

Create another table.

Create another table.

Pets			
PetID	PetName	PetDOB	FriendID
1	Chikin	24/09/2016	2
2	Cauchy	01/03/2012	3
3	Gauss	01/03/2012	3

Create another table.

Pets			
PetID	PetName	PetDOB	FriendID
1	Chikin	24/09/2016	2
2	Cauchy	01/03/2012	3
3	Gauss	01/03/2012	3

Foreign key

>>> The foreign key 'points at' the primary key

Pets

PetID	PetName	…	FriendID
1	Chikin	…	2
2	Cauchy	…	3
3	Gauss	…	3

Friends

FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Pets			
PetID	PetName	…	FriendID
1	Chikin	…	2
2	Cauchy	…	3
3	Gauss	…	3

Many

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

>>> The foreign key 'points at' the primary key
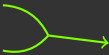
Pets

PetID	PetName	…	FriendID
1	Chikin	…	2
2	Cauchy	…	3
3	Gauss	…	3

Friends

FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

```
>>> Check that we fixed all these problems



  * Have to store NULL in every entry with no pet
  * What if I meet a friend with 3+ pets? Many more NULLs
  * New one-to-many relationship between pets and toys?
  * Pets are tied to owners. Delete an owner → delete pets
  * Ambiguity. Is information related to pets or owners?
```

FriendsPets						
PetID	PetName	…	FriendID	FriendID	FirstName	…
1	Chikin	…	2	2	Y	…
2	Cauchy	…	3	3	Z	…
3	Gauss	…	3	3	Z	…

FriendsPets						
PetID	PetName	…	FriendID	FriendID	FirstName	…
1	Chikin	…	2	2	Y	…
2	Cauchy	…	3	3	Z	…
3	Gauss	…	3	3	Z	…

Primary/foreign key pair

Challenge: Can you create a one-to-many relationship between Friends and Friends? How will you model it?

* Game in which friends fight to the death. A friend can
  beat many others, but can only be beaten by one at most.

Friends				
FriendID	FirstName	LastName	FavColour	DefeatedByID
1	X	A	red	2
2	Y	B	blue	NULL
3	Z	C	NULL	2

* Foreign key 'points at' the primary key
* Two rows <u>can</u> share same foreign key value
* Two rows <u>can not</u> share same primary key value
* Primary key can never be NULL
* All tables <u>should</u> have a primary key
* A PK or FK can be made of more than one column.

* Foreign key 'points at' the primary key
* Two rows <u>can</u> share same foreign key value
* Two rows <u>can not</u> share same primary key value
* Primary key can never be NULL
* All tables <u>should</u> have a primary key
* A PK or FK can be made of more than one column.

For example, a company might sell group holiday packages and
the primary key of their Customer table might be made of a
GroupID and GroupMemberNumber.

When there is a foreign key entry that is not NULL, the
primary key entry that it 'points at' <u>must</u> exist.

When there is a foreign key entry that is not NULL, the primary key entry that it 'points at' <u>must</u> exist.

Guarantees that a foreign key is not 'meaningless'.

Guarantees that a foreign key is not 'meaningless'.

Friends				
FriendID	FirstName	LastName	FavColour	DefeatedByID
1	X	A	red	4
2	Y	B	blue	NULL
3	Z	C	NULL	2

>>> Identifying a primary / foreign key pair?

A foreign key is any column (or collection of columns) where each record is guaranteed to equal one, and only one, primary key entry in the other table.

Problem: What happens if the database is sloppy, and there aren't any foreign keys??

```
Houses(Bedrooms, Bathrooms, LandSize, PostCode)
Suburbs(PostCode, SuburbName).
```

```
        Houses(Bedrooms, Bathrooms, LandSize, PostCode)
        Suburbs(PostCode, SuburbName).
```

1. Is every PostCode entry in Suburbs unique?
2. Is every PostCode in Houses also in Suburbs?

Does it really matter if we can't tell?

```
Houses(Bedrooms, Bathrooms, LandSize, PostCode)
Suburbs(PostCode, SuburbName).
```

1. Is every PostCode entry in Suburbs unique?
2. Is every PostCode in Houses also in Suburbs?

Does it really matter if we can't tell?

  * From 1: can't be sure which suburb a house is in.
  * From 1: joining can lead to unexpected duplicates.
  * From 2: can't find any matching suburb.

* A class has many students,
  and a student attends many classes.
* A company has many investors,
  and an investor invests in many companies.
* A person engages with many government departments,
  and a government department engages with many people.

* Each friend can scratch many backs,
  and a back can be scratched by many friends

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Scratched			
ScratcherID	Date	Time	ScratcheeID
1	05/09/2018	12:00pm	2
1	05/09/2018	12:30pm	3
2	06/09/2018	11:00am	1
3	07/09/2018	10:00am	1

  * Each friend can scratch many backs,
    and a back can be scratched by many friends

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Scratched			
ScratcherID	Date	Time	ScratcheeID
1	05/09/2018	12:00pm	2
1	05/09/2018	12:30pm	3
2	06/09/2018	11:00am	1
3	07/09/2018	10:00am	1

>>> Many-to-many relationship

* Each friend can scratch many backs,
  and a back can be scratched by many friends

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Scratched			
ScratcherID	Date	Time	ScratcheeID
1	05/09/2018	12:00pm	2
1	05/09/2018	12:30pm	3
2	06/09/2018	11:00am	1
3	07/09/2018	10:00am	1

* Each friend can scratch many backs,
  and a back can be scratched by many friends

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Scratched			
ScratcherID	Date	Time	ScratcheeID
1	05/09/2018	12:00pm	2
1	05/09/2018	12:30pm	3
2	06/09/2018	11:00am	1
3	07/09/2018	10:00am	1

* Each friend can scratch many backs,
  and a back can be scratched by many friends

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Friends		
FriendID	FirstName	…
1	X	…
2	Y	…
3	Z	…

Scratched			
ScratcherID	Date	Time	ScratcheeID
1	05/09/2018	12:00pm	2
1	05/09/2018	12:30pm	3
2	06/09/2018	11:00am	1
3	07/09/2018	10:00am	1

Friend_Scratched_Friend								
FrID	FriendName	…	SrID	…	SeID	FrID	FriendName	…
1	X	…	1	…	2	2	Y	…
1	X	…	1	…	3	3	Z	…
2	Y	…	2	…	1	1	X	…
3	Z	…	3	…	1	1	X	…

Friend_Scratched_Friend								
FrID	FriendName	…	SrID	…	SeID	FrID	FriendName	…
1	X	…	1	…	2	2	Y	…
1	X	…	1	…	3	3	Z	…
2	Y	…	2	…	1	1	X	…
3	Z	…	3	…	1	1	X	…

Pair 1

Friend_Scratched_Friend								
FrID	FriendName	…	SrID	…	SeID	FrID	FriendName	…
1	X	…	1	…	2	2	Y	…
1	X	…	1	…	3	3	Z	…
2	Y	…	2	…	1	1	X	…
3	Z	…	3	…	1	1	X	…

Pair 2

 * A friend can play with many pets,
   and a pet can play with many friends

Pets		
PetID	PetName	…
1	Chikin	
2	Cauchy	
3	Gauss	

Friends		
FriendID	FirstName	…
1	X	
2	Y	
3	Z	

PlayCount		
PetID	Count	FriendID
1	3	1
1	5	2
3	4	2

* A person can have at most one head,
  and each head belongs to only one person
* A table record has exactly one primary key value,
  and each primary key value belongs to exactly one record
* A user has one set of log-in details,
  and each set of log-in details belong to one user

* One friend can have at most one passport, and each passport belongs to only one friend

Friends					
FriendID	FirstName	…	PptCountry	PptNo	PptExpiry
1	X		Australia	E1321	12/03/2021
2	Y		New Zealand	LA123	01/09/2032
3	Z		Monaco	S9876	19/06/2028

```
>>> Why not keep one-to-one relationships in the same table?
```

  * NULLs (many passport attributes? few people have them?)

```
>>> Why not keep one-to-one relationships in the same table?




   * NULLs (many passport attributes? few people have them?)
   * Dependence: Delete friend → delete passport.
```

```
>>> Goodbye, Mr. X
```

Friends					
FriendID	FirstName	…	PptCountry	PptNo	PptExpiry
2	Y		New Zealand	LA123	01/09/2032
3	Z		Monaco	S9876	19/06/2028

How do we delete a friend without deleting their passport?

How do we delete a friend without deleting their passport?

Passports			
PptNo	PptCountry	PptExpiry	FriendID
E1321	Australia	12/03/2021	NULL
LA123	New Zealand	01/09/2032	2
S9876	Monaco	19/06/2028	3

How do we delete a friend without deleting their passport?

Passports			
PptNo	PptCountry	PptExpiry	FriendID
E1321	Australia	12/03/2021	NULL
LA123	New Zealand	01/09/2032	2
S9876	Monaco	19/06/2028	3

Mr. X

>>> Any problems with this approach though?

How do we delete a friend without deleting their passport?

Passports			
PptNo	PptCountry	PptExpiry	FriendID
E1321	Australia	12/03/2021	NULL
LA123	New Zealand	01/09/2032	2
S9876	Monaco	19/06/2028	3

```
>>> Any problems with this approach though?
```

How do we delete a friend without deleting their passport?

Passports			
PptNo	PptCountry	PptExpiry	FriendID
E1321	Australia	12/03/2021	NULL
LA123	New Zealand	01/09/2032	2
S9876	Monaco	19/06/2028	3

Deleting a friend will delete the owner's name

We should avoiding keeping the person's name in both tables, since otherwise we have redundant data.

`>>>` Any idea how to fix this?

  * Create 'people' table with binary variable for friend?

Definition
A binary variable is always either 0, 1 or NULL.
Usually, 0 represents false and 1 repesents true.

`>>> Any idea how to fix this?`

* Create 'people' table with binary variable for friend?
* Create separate tables for friends, enemies, etc…?

    Leave it to the database designers.

  * Missing information
  * Conflicting information (due to redundancy)
  * Not enough levels of a categorical variable
  * Binary answer when binary is not appropriate
  * Hard to join the tables and connect records
  * Hard to search for information in the database
  * Many more, keep eyes open…

Enjoy.

```
>>> Where are we now?
```

Day 1
 1. Introduction
 2. Intro to relational model
 3. Tables and relationships
 4. Programming in SQL
 5. Basic SQL
 6. Joining in SQL

Page is hyperlinked: click a topic above to jump to it.

Time for the real deal

If you've completed set-up, that's great!
Otherwise, we can troubleshoot it this afternoon.

If you're on macOS, don't worry!
I'll be giving a demo of Sequel Ace later.

In Azure Data Studio, I'll do the following:

  * Connect to 'localhost'.
  * Open a new query tab.
  * Change between databases.
  * Figure out what tables are in a database.
  * Explain what a schema is.
  * Figure out what columns are in a table.
  * Figure out what the data types are.
  * Figure out what the primary/foreign keys are.
  * Figure out if NULL values are allowed.

>>> Bonus demo

Sneak preview of SQL code

* The USE clause.
* Retrieve Friends.
* Retrieve Pets.
* Join Friends with Pets.
* Aliases.
* Quoting identifiers.

SeLeCt*FrOm[NoTeS].
[pEtS]rIpHaRaMbE20160528

---

* Upper/lower-case has no effect
* Spaces usually have no effect
* Square brackets can be omitted
* New lines have no effect
* Alias can be almost anything

                So pay attention to style

The concept of an alias is explained on the next slide.

Aliases give temporary names to tables, and should be used to simplify and shorten your queries.

Without aliases:

```
SELECT *
FROM Notes.Friends JOIN Notes.Pets
ON Notes.Friends.friendID = Notes.Pets.friendID;
```

With aliases:

```
SELECT *
FROM Notes.Friends F JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

From now on, we will **always use aliases**.

```
>>> A note on syntax




Another (optional) way to write aliases
_____

SELECT *
FROM Notes.Friends AS F JOIN Notes.Pets AS P
ON F.friendID = P.friendID;
_____
```

```
>>> Where are we now?


Day 1
 1. Introduction
 2. Intro to relational model
 3. Tables and relationships
 4. Programming in SQL
 5. Basic SQL
 6. Joining in SQL


Page is hyperlinked: click a topic above to jump to it.
```

>>> Time to learn SQL

We will now look closely at some basic clauses

```
>>> SQL clause: FROM


The FROM clause specifies table(s)
to access in the SELECT statement (and others).

────────────────────────────────────────────────────────────

FROM MySchema.MyTable MyAlias

────────────────────────────────────────────────────────────

The above will not run because there is no SELECT. You'll use
FROM in almost every query, though.
```

The FROM clause specifies table(s)
to access in the SELECT statement (and others).

FROM MySchema.MyTable MyAlias

The above will not run because there is no SELECT. You'll use
FROM in almost every query, though.

Remember: MySQL doesn't have schemas (but don't Google it)

The SELECT clause allows you to choose columns.
You can select all columns with SELECT *

We will look at the execution of this query:

```sql
SELECT F.FirstName, F.FavColour
FROM Notes.Friends F;
```

Note: the alias F seems to have been used before it was
created! We will learn about the (sometimes confusing) SQL
order of execution.

FROM Notes.Friends

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

SELECT F.FirstName, F.FavColour

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

result

Unnamed	
FirstName	FavColour
X	red
Y	blue
Z	NULL

But did you see that order of execution?

But did you see that order of execution?

* Syntactic order of execution
* Logical order of execution
* Optimal order of execution

>>> SQL clause: WHERE

The WHERE clause allows you to choose rows,
using a search condition.

We will look at the execution of this query:

```sql
SELECT F.firstName, F.lastName
FROM Notes.Friends F
WHERE favColour = 'red';
```

FROM Notes.Friends

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

WHERE FavColour = 'red'

Friends			
FriendID	FirstName	LastName	FavColour
1	X	A	red
2	Y	B	blue
3	Z	C	NULL

SELECT FirstName, LastName

Unnamed			
ID	FirstName	LastName	FavColour
1	X	A	red

result

Unnamed	
FirstName	LastName
X	A

1. FROM
2. WHERE
3. SELECT

```
>>> Order of execution
```

Order of execution is irrelevant, Danny!

>>> Order of execution

Order of execution is irrelevant, Danny!

Wrong you are.

Aliases can be created in the SELECT clause too!

```sql
SELECT F.FirstName AS Nombre, F.FavColour AS ColorFavorito
FROM Notes.Friends F
WHERE ColorFavorito = 'red';
```

Let's try executing the above. What will happen?

```
>>> Why do you keep saying 'clause'?
```

SQL is like speaking ... or cooking.

* Clauses are components of statements.
* The statements we're learning are called queries.
* A statement is somewhat comparable to a 'sentence'.
* Better to think of them as ingredients in a recipe?

* We've seen how to 'chop' (with SELECT and WHERE).
* We've seen how to 'change' (with table/column aliases).

Can we also change the entries?

* We've seen how to 'chop' (with SELECT and WHERE).
* We've seen how to 'change' (with table/column aliases).

Can we also change the entries?

Change entries with the CASE WHEN expression.

```
SELECT *, CASE WHEN FavColour = 'red' THEN 'rojo'
              WHEN FavColour = 'blue' THEN 'azul'
              ELSE FavColour END AS ColorFavorito
FROM Notes.Friends;
```

Let's execute the above. What will it do?

We can also reorder the results!

```sql
SELECT *
FROM Notes.Friends
ORDER BY FriendID DESC;
```

Let's execute it to experiment.

>>> Lexicographic ordering

What happens if we order by a character string?

Numbers	
Num	NumString
111	'111'
31	'31'
32	'32'
211	'211'

```
SELECT *
FROM Notes.Numbers
ORDER BY NumString;
```

Let's execute it to find out.

Functions that transform entries are often called scalar functions. Perhaps the most important is:

```
SELECT *
FROM Notes.Numbers
ORDER BY CAST( NumString AS INT );
```

Let's execute it.

>>> Other scalar functions

The three categories we will look at are

  * Mathematical functions
  * String functions
  * Date and time functions

Any many more (click here).

Function	Description
SQRT	Square root
ROUND	Rounding
RAND	Generate random number

I will go through some examples...

```
>>> Two string functions
```

Function	Description
CONCAT	Concatenate columns
SUBSTRING	Extract characters

I will go through some examples...

Function	Description
DAY	Extract the day (of the month)
MONTH	Extract the month
YEAR	Extract the year

I will go through some examples...

```
>>> Where are we now?


Day 1
 1. Introduction
 2. Intro to relational model
 3. Tables and relationships
 4. Programming in SQL
 5. Basic SQL
 6. Joining in SQL


Page is hyperlinked: click a topic above to jump to it.
```

A JOIN (also known as an INNER JOIN) pairs the records from one table with the records from another table, using a primary/foreign key pair.

We will look at the execution of this query:

```
SELECT F.firstName, P.petName
FROM Notes.Friends F JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

We will look at the execution of this query:

```sql
SELECT F.firstName, P.petName
FROM Notes.Friends F JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

Another way to write the same query: implicit syntax

```sql
SELECT F.firstName, P.petName
FROM Notes.Friends F, Notes.Pets P
WHERE F.friendID = P.friendID;
```

Yet another way to write the same query:

```sql
SELECT F.firstName, P.petName
FROM Notes.Friends F INNER JOIN Notes.Friends P
ON F.friendID = P.friendID
```

```
>>> SQL query: JOIN
```

Note that JOIN is an operator that is inside the FROM clause.

```
FROM Friends F JOIN Pets P ON F.FriendID = P.FriendID
```

Pets			
PetID	PetName	…	FriendID
1	Chikin		2
2	Cauchy		3
3	Gauss		3

Friends		
FriendID	FirstName	…
1	X	
2	Y	
3	Z	

SELECT F.FirstName, P.PetName

Unnamed						
PetID	PetName	…	FriendID	FriendID	FirstName	…
1	Chikin	…	2	2	Y	…
2	Cauchy	…	3	3	Z	…
3	Gauss	…	3	3	Z	…

result

Unnamed	
PetName	FirstName
Chikin	Y
Cauchy	Z
Gauss	Z

JOIN is technically an operator, not a clause.

1. FROM   (and JOIN)
2. WHERE
3. SELECT

Table1		
A	B	C
1	Ignorance	is
2	War	is
3	Freedom	is
4	Friendship	is

Table2		
D	E	A
slavery.	3	1
weakness.	4	2
strength.	1	3
peace.	2	4

* SELECT B,C,D FROM Table1 T1, Table2 T2 WHERE T1.A = T2.A
* SELECT B,C,D FROM Table1 T1, Table2 T2 WHERE T1.A = T2.E

```
>>> Solutions

  * SELECT B,C,D FROM Table1 T1, Table2 T2 WHERE T1.A = T2.A
```

B	C	D
Ignorance	is	slavery.
War	is	weakness.
Freedom	is	strength.
Friendship	is	peace.

* SELECT B,C,D FROM Table1 T1, Table2 T2 WHERE T1.A = T2.A

B	C	D
Ignorance	is	slavery.
War	is	weakness.
Freedom	is	strength.
Friendship	is	peace.

* SELECT B,C,D FROM Table1 T1, Table2 T2 WHERE T1.A = T2.E

B	C	D
Ignorance	is	strength.
War	is	peace.
Freedom	is	slavery.
Friendship	is	weakness.

The join query below (that we looked at earlier)
excludes any friends that have no pets (and vice versa).

```
SELECT F.firstName, P.petName
FROM Notes.Friends F JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

The join query below (that we looked at earlier)
excludes any friends that have no pets (and vice versa).

```sql
SELECT F.firstName, P.petName
FROM Notes.Friends F JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

LEFT JOIN keeps every row from the table on the left.

```sql
SELECT F.firstName, P.petName
FROM Notes.Friends F LEFT JOIN Notes.Pets P
ON F.friendID = P.friendID;
```

```
>>> SQL query: LEFT JOIN. Remember this?
```

FROM Friends F JOIN Pets P ON F.FriendID = P.FriendID

Pets			
PetID	PetName	…	FriendID
1	Chikin		2
2	Cauchy		3
3	Gauss		3

Friends		
FriendID	FirstName	…
1	X	
2	Y	
3	Z	

```
>>> The result was...
```

Unnamed						
PetID	PetName	…	FriendID	FriendID	FirstName	…
1	Chikin	…	2	2	Y	…
2	Cauchy	…	3	3	Z	…
3	Gauss	…	3	3	Z	…

If we did a LEFT JOIN instead we would get:

FROM Friends F LEFT JOIN Pets P ON F.FriendID = P.FriendID

Unnamed					
PetID	PetName	…	FriendID	FriendID	FirstName	…
NULL	NULL	…	NULL	1	X	…
1	Chikin	…	2	2	Y	…
2	Cauchy	…	3	3	Z	…
3	Gauss	…	3	3	Z	…

```
>>> SQL query: LEFT JOIN
```

result

Unnamed	
PetName	FirstName
NULL	X
Chikin	Y
Cauchy	Z
Gauss	Z

Question for the class:
What does RIGHT JOIN do?

Do exercises at the ends of Chapters 1 and 2.

Exercises 2.5.5, 2.5.6, 2.7.3 and 2.7.4 all use some material introduced in tomorrow's session.

Click here to open the textbook.