

Unleashing SQL

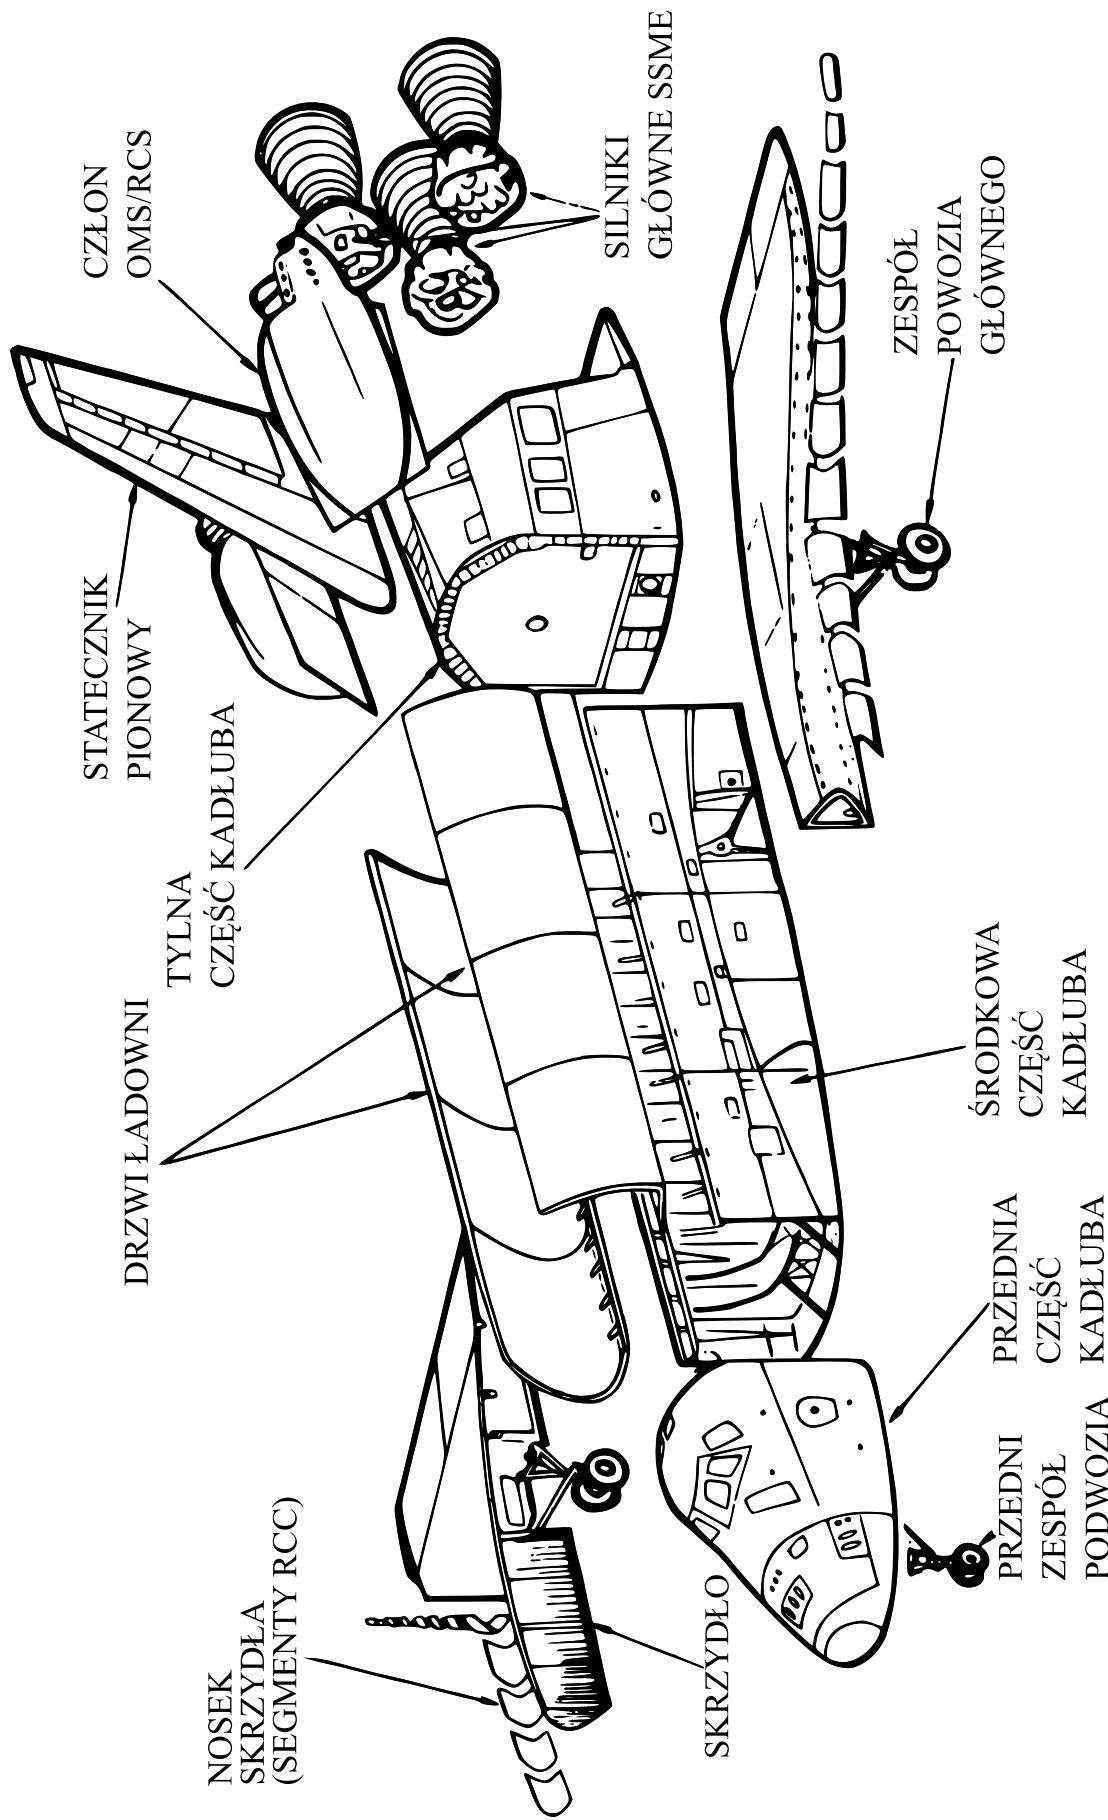


with R

Daniel Fryer

NZSSN

2022/02/18



The fundamentals

Why learn R?

Community!

- Nerds
- Statisticians
- Data analysts
- Social scientists
- Psychologists
- Bioinformaticians
- Medical researchers
- Programmers
- Artists
- Machine learners
- Industry
- ...

Tons of help and inspiration:

- Twitter
- Blogs
- Educators
- Books
- Galleries
- Journals
- Conferences
- CRAN

Powerful and capable:

- Reproducible research!
- Make slides ([xaringan](#))
- Automatic reports ([RMarkdown](#))
- Dashboards ([Shiny](#))
- Interactive plots ([plotly](#))
- Used in industry

State of the art statistical analysis
packages.

Free, open source, transparent.

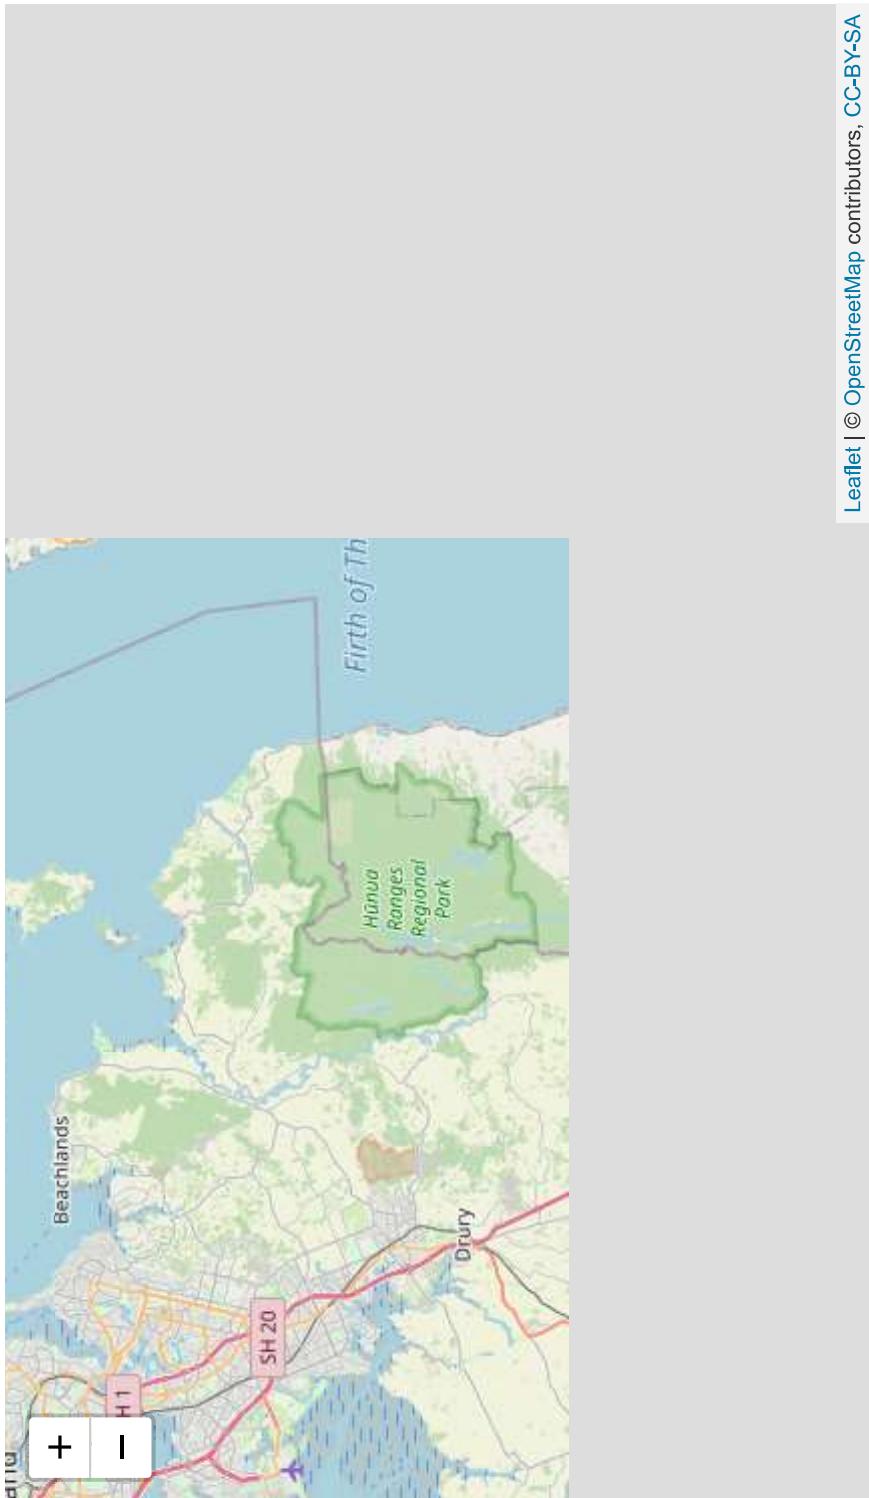
It's from New Zealand!

Why learn R with SQL?

- Keep your data organised.
- Understand **tidy data** better.
- Understand the **tidyverse** better.
- Work with remote servers in large collaborations.
- Combine all your CSV files into a single file (**SQLite**).
- Search and query your combined CSV files with ease (**DB Browser**).

Showing off

```
library(leaflet)  
leaflet() %>% addTiles() %>% setView(174.76898, -36.85231, zoom = 10)
```



Project management



Scripts are essentially text files that you save your code to.

Every **project** should have its own folder (usually with subfolders).

A single project typically involves:

- Preparation, analysis, presentation, data.

The **working directory** should always be the **project root**.

Your **workspace** is the RStudio environment, including all the variables you have created, and your 'history'. *Do not rely on this.*

"The source code is real. The objects are realizations of the source code."

Where do you put your data?

In the data folder!

- Raw data
- Processed data

umbrella

 | README.md

 | start-here.R

 | data

 | raw

 | processed

Always think about a new person opening your project directory for the first time. Will it work seamlessly? Will they know how to use it?

 | prepare

 | cleaning-functions.R

Restart your session frequently.

 | analyse

 | analysis-functions.R

 | present

 | figures

 | unleashing-SQL

Umbrella: a simple example project

[Click here to open](#)

R is a calculator

```
# a boring calculation  
5 + 7
```

```
# [1] 12
```

A calculator that can save variables

```
x <- 5 + 7  
x*2
```

```
## [1] 24
```

In all sorts of ways

```
x <- c(1, 2, 3, 4, 5)  
x*2
```

```
## [1] 2 4 6 8 10
```

R has lists

These are very powerful

```
Friends <- list(  
  id = c(1, 2, 3),  
  FirstName = c('X', 'Y', 'Z'),  
  LastName = c('A', 'B', 'C'),  
  FavColour = c('red', 'blue', NA)  
)  
Friends  
  
## $id  
## [1] 1 2 3  
## $FirstName  
## [1] "X" "Y" "Z"  
## $LastName  
## [1] "A" "B" "C"  
## $FavColour  
## [1] "red" "blue" NA
```

But the output looks kind of weird.

Let's make the lists look better

```
Friends <- data.frame(Friends  
Friends
```

```
##   id FirstName LastName FavColour  
## 1  1        X       A     red  
## 2  2        Y       B    blue  
## 3  3        Z       C    <NA>
```

🤔 Ummmm... can we do better?

```
library(tibble)  
Friends <- tibble(Friends  
Friends
```

```
## # A tibble: 3 × 4  
##   id FirstName LastName FavColour  
##   <dbl> <chr>    <chr>    <chr>  
## 1     1      X       A     red  
## 2     2      Y       B    blue  
## 3     3      Z       C    <NA>
```

🧠 Wow, that's informative!

But *can we do better?*

Showing 1 to 3 of 3 entries			
	id	FirstName	LastName
			FavColour
1	1	X	A
2	2	Y	B
3	3	Z	C

Search:

Show 10 entries

Previous Next

A better dataset

Show 8 ✓ entries

Search:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
9	4.4	2.9	1.4	0.2	setosa
12	4.8	3.4	1.6	0.2	setosa
15	5.8	4	1.2	0.2	setosa
17	5.4	3.9	1.3	0.4	setosa
18	5.1	3.5	1.4	0.3	setosa

Showing 1 to 8 of 50 entries

Previous Next

Yet underneath, 'tis still but a list

```
str(Friends)
```

```
### tibble [3 x 4] (S3:tbl_df/tbl/data.frame)
### $ id      : num [1:3] 1 2 3
### $ FirstName: chr [1:3] "X" "Y" "Z"
### $ LastName : chr [1:3] "A" "B" "C"
### $ FavColour: chr [1:3] "red" "blue" NA
```

Lists, precious lists

-  SQL stores data in a way that is great for machines
-  R stores data in a way that is great for programmers

Everything is...

Everything in R is either **data** or a **function**. We refer to 'things' as 'objects'.

The contents of lists can be accessed with \$

```
Friends$FavColour
```

```
## [1] "red"  "blue" NA
```

So, what is \$?

Well, if \$ is not data, then it's a function.

```
`$`(Friends, FavColour)
```

```
## [1] "red"  "blue" NA
```

Just a sneaky function.

Here's another function for accessing things! [

```
Friends[1,4]
```

```
## # A tibble: 1 × 1
##   FavColour
##   <chr>
## 1 red
```

Thinking in vectors

Here's a vector of TRUE / FALSE

```
x <- c(TRUE, FALSE, TRUE)
```

We can use it to get friends.

```
Friends[x, ]
```

```
## # A tibble: 2 x 4
##   id FirstName LastName FavColour
##   <dbl> <chr>    <chr>   <chr>
## 1     1 X          A        red
## 2     2 Z          C        <NA>
```

```
Friends[!x, ]
```

```
## # A tibble: 1 x 4
##   id FirstName LastName FavColour
##   <dbl> <chr>    <chr>   <chr>
## 1     1 Y          B        blue
```

Packages make functions and data

But we can (and should) make our own functions.

```
# A function that checks if x is 'blue'  
is_x_blue <- function(x) {  
  x == 'blue'  
}
```

And we should use them often.

```
is_x_blue('green')
```

```
## [1] FALSE
```

```
is_x_blue('blue')
```

```
## [1] TRUE
```

We can use the argument explicitly, if we like.

```
is_x_blue(x='red')
```

```
## [1] FALSE
```

Functions are building blocks.

The longer code gets, the harder it is to think about.

Do you prefer this?

```
spec <- iris$Species  
iris_setosa <- iris[spec == "setosa", ]  
m <- lm(Sepal.Length ~ Petal.Length, data=iris_setosa)  
coef(summary(m))
```

Or this?

```
# Get the iris Setosa species  
iris_setosa <- where_species_is_setosa(iris)  
  
# Fit Model (1), and get results  
results <- fit_linear_model(iris_setosa, number = 1)  
  
# Print results  
results
```

The pipe operator %>%

```
# Get the iris Setosa species
iris_setosa <- where_species_is_setosa(iris)

# Fit Model (1), and get results
results <- fit_linear_model(iris_setosa, number = 1)

# Print results
results
```

Many people prefer this instead:

```
# Get Setosa species and fit Model (1)
iris %>%
  where_species_is_setosa() %>%
  fit_linear_model(number = 1)
```

The pipe comes from  magrittr

```
library(magrittr)
```



Tidy data

*"Happy families are all alike;
every unhappy family is unhappy in its own way"*

- Leo Tolstoy.

*"Tidy datasets are all alike,
but every messy dataset is messy in its own way"*

- Hadley Wickham.

*"The principles of tidy data are closely tied to those of relational databases... but
are framed in a language familiar to statisticians"* [1]

[1] Wickham, H (2014), *Tidy Data*, The R Journal

Tidy data principles

1. Every variable is a column.
2. Every observation is a row.
3. Every cell holds a single (atomic) value.

Conditions 1 and 2 are often referred to as **long format**.

Advantages

- Datasets can be read and understood universally.
- Easier for R packages to work together.
- Enables SQL-like operations.



Messy data

The R package  `tidyverse` has functions for tidying messy data.

And it also has some datasets for us to play with.

Show 3 ▾ entries							Search: <input type="text"/>
	religion	<\$10k	\$10-20k	\$20-30k	\$30-40k	\$40-50k	
1	Agnostic	27	34	60	81	76	
2	Atheist	12	27	37	52	35	
3	Buddhist	27	21	30	34	33	

Showing 1 to 3 of 18 entries

Previous [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) Next



Cleaning religion_income

This kind of dataset is sometimes referred to as **wide format**.

tidyR gives us `pivot_longer`

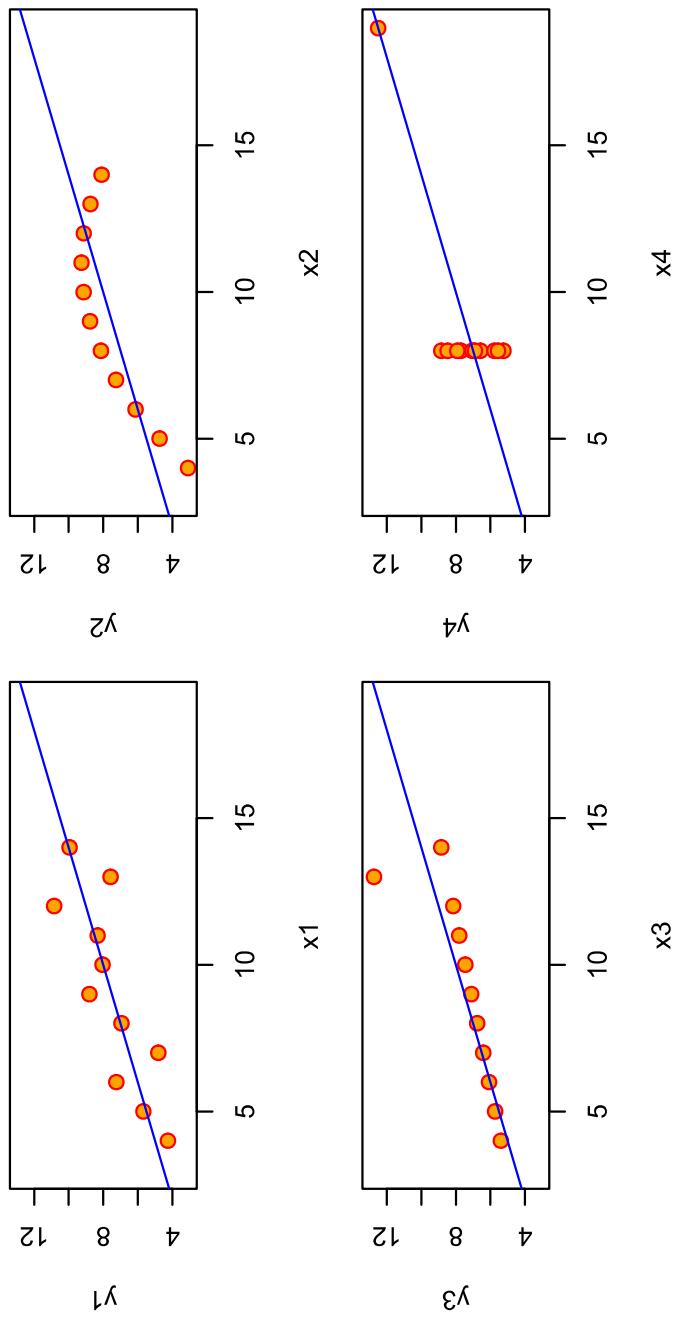
```
pivot_longer(  
  data = relig_income,  
  cols = !religion,  
  names_to = "income",  
  values_to = "count"  
)
```

Show	3	▼ entries	Search: <input type="text"/>	religion	income	count
1	Agnostic	<\$10k	27			
2	Agnostic	\$10-20k	34			
3	Agnostic	\$20-30k	60			

Showing 1 to 3 of 180 entries



Messy data no.2: Anscombe's Quartet



Anscombe's Quartet

Show 6 ✓ entries

	x1 ◀	x2 ◀	x3 ◀	x4 ◀	y1 ◀	y2 ◀	y3 ◀	y4 ◀
1	10	10	10	8	8.04	9.14	7.46	6.58
2	8	8	8	8	6.95	8.14	6.77	5.76
3	13	13	13	8	7.58	8.74	12.74	7.71
4	9	9	9	8	8.81	8.77	7.11	8.84
5	11	11	11	8	8.33	9.26	7.81	8.47
6	14	14	14	8	9.96	8.1	8.84	7.04

Showing 1 to 6 of 11 entries

Previous 1 2 Next

Search:



Cleaning Anscombe's Quartet

```
pivot_longer(  
  data = anscombe,  
  cols = everything(),  
  names_to = c("value", "set"),  
  names_pattern = "(.)(.)"  
)
```

Show 4 ✓ entries

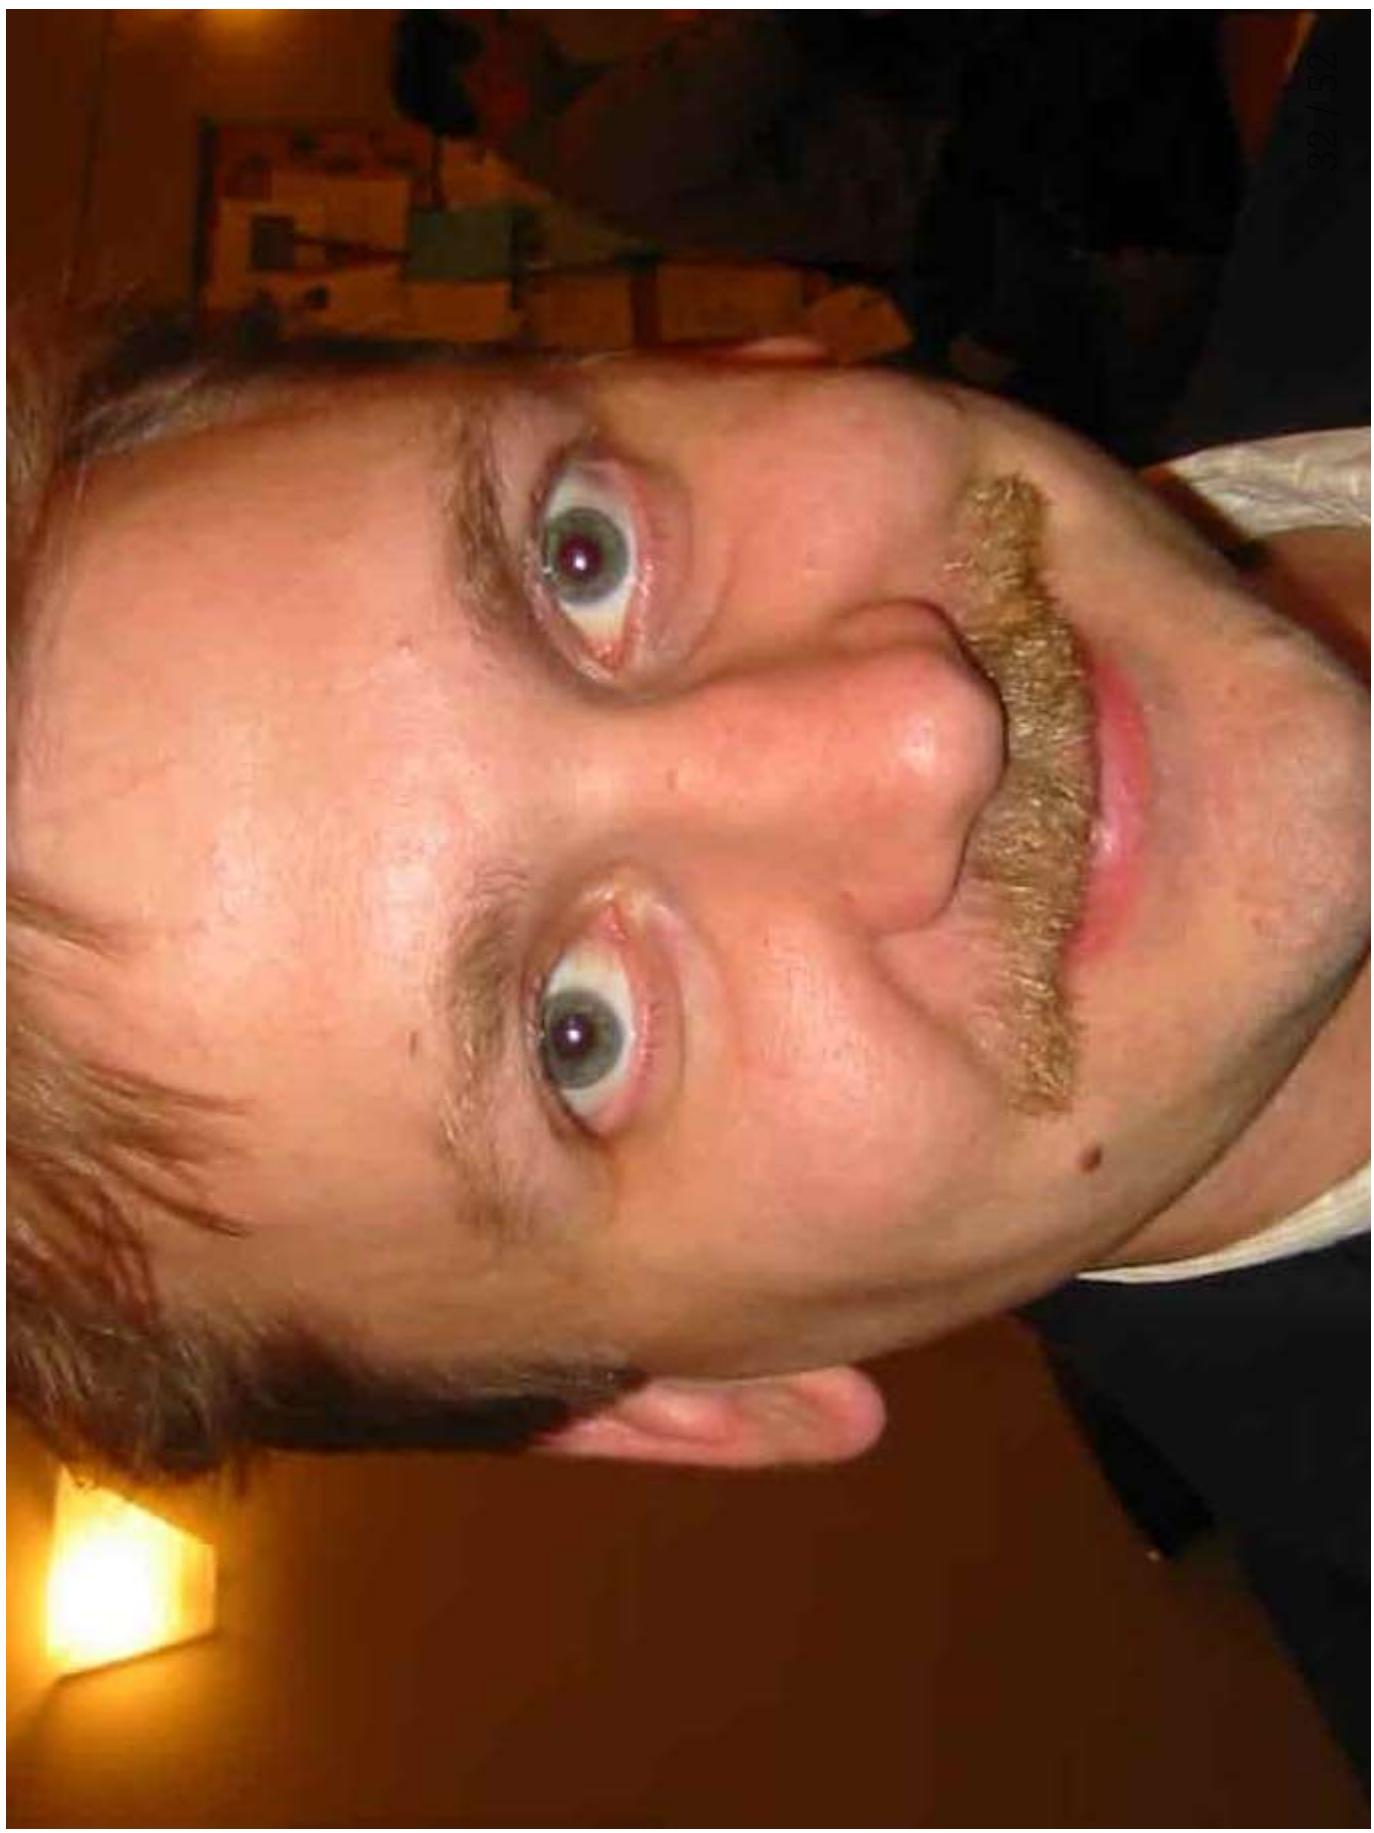
	set	x	y
1	1	10	8.04
2	2	10	9.14
3	3	10	7.46
4	4	8	6.58

Search:

Showing 1 to 4 of 44 entries

Previous 1 2 3 4 5 ... 11 Next





32 / 52

Getting connected to SQL

*For MySQL and T-SQL, see
[this guide](#)
covering local and remote connections from R*

The wonderful SQLite

SQLite is:

- Small
- Fast
- Self contained
- High reliability
- Full-featured

The most used database engine in the world.

Around one trillion active databases in use.

Hundreds of SQLite databases on any smart phone.

One of the top 5 most widely deployed software modules of any kind.



Connect or create, one line!

First get our directories organised (see [umbrella](#)).

```
library(here)
```

Now, create or connect

```
library(RSQLite)
con <- dbConnect(
  SQLite(),
  here("data", "raw", "Sandpit.sqlite")
)
```

If it can't be found, then Sandpit.sqlite will be created (empty).

Use the Sandpit database

The package  `dplyr` gives us a whole 'grammar of data manipulation'.

The package  `dbplyr` allows `dplyr` to talk to SQL.

```
library(dplyr)  
library(dbplyr)
```

So many packages! From now on I'll write them like this:

```
dplyr::tbl()
```

The `dplyr::` means we are using a function from `dplyr`. The function we're using is called `tbl`.



Connect to a table

```
banana <- dplyr::tbl(con, "Ape_Banana")  
banana
```

```
## # Source:   table<Ape_Banana> [? x 7]  
## # Database: SQLite 3.36.0  
## # [D:\CloudDrive\OneDrive\Cloud-drive\Teach\Courses\Workshops\SQL\repositories\umbrella  
## # BananaID TasteRank DatePicked DateEaten Ripe TreeID Comments  
## # <int> <int> <int> <int> <int> <chr>  
## # 1       1       2       20181003 20181004 0  1 <NA>  
## # 2       2       4       20181003 20181004 1  2 <NA>  
## # 3       3       4       20181003 20181004 1  2 <NA>  
## # 4       4       5       20181003 20181006 1  1 <NA>  
## # 5       5       5       20181003 20181006 1  2 best banana ever  
## # 6       6       3       20181003 20181004 1  2 <NA>  
## # 7       7       2       20181002 20181004 0  3 <NA>  
## # 8       8       5       20181002 20181005 1  3 smooth and delectable  
## # 9       9       3       20181002 20181003 1  4 <NA>  
## # 10      10      3       20181002 20181003 1  5 <NA>  
## # ... with more rows
```

Grammar of data manipulation SQL

The function `dplyr::filter` is like the SQL WHERE clause.

```
ripe_banana <- banana %>% dplyr::filter(Ripe == 1)  
ripe_banana
```

```
## # Source: lazy query [?? x 7]  
## # Database: sqlite 3.36.0  
## # [D:\CloudDrive\OneDrive\Cloud-drive\Teach\Courses\Workshops\SQL\repositories\umbrella  
## # BananaID TasteRank DatePicked DateEaten Ripe TreeID Comments  
## # <int> <int> <int> <int> <int> <int> <chr>  
## 1 2 4 20181003 20181004 1 2 <NA>  
## 2 3 4 20181003 20181004 1 2 <NA>  
## 3 4 5 20181003 20181006 1 1 <NA>  
## 4 5 5 20181003 20181006 1 2 best banana ever  
## 5 6 3 20181003 20181004 1 2 <NA>  
## 6 8 5 20181002 20181005 1 3 smooth and delectable  
## 7 9 3 20181002 20181003 1 4 <NA>  
## 8 10 3 20181002 20181003 1 5 <NA>  
## 9 12 5 20181002 20181005 1 4 <NA>  
## 10 16 5 20181001 20181004 1 5 a culinary delight  
## # ... with more rows
```

Grammar of data manipulation SQL

But seriously, dplyr::filter is *actually* the WHERE clause.

```
ripe_banana %>% dplyr::show_query()
```

```
## #<SQL>
## SELECT *
## FROM `APE_Banana`
## WHERE (`Ripe` = 1.0)
```

To execute the query:

```
ripe_banana %>% dplyr::collect()
```

```
## # A tibble: 34 x 7
##   BananaID TasterRank DatePicked DateEaten Ripe TreeID Comments
##       <int>      <int>      <int>      <int> <int> <chr>
## 1           2          4        20181003  20181004     1    2 <NA>
## 2           3          4        20181003  20181004     1    2 <NA>
## 3           4          5        20181003  20181006     1    1 <NA>
## 4           5          5        20181003  20181006     1    2 best banana ever
## 5           6          3        20181003  20181004     1    2 <NA>
## 6           8          5        20181002  20181005     1    3 smooth and delectable
## 7           9          3        20181002  20181003     1    4 <NA>
## 8          10          3        20181002  20181003     1    5 <NA>
## 9          12          5        20181002  20181005     1    4 <NA>
```



We can write our own SQL

```
DBI::dbGetQuery(con,  
"SELECT TreeID, COUNT(*) AS NumRipe, AVG(TasteRank) AS AvgTaste  
FROM Ape_Banana  
WHERE DatePicked > '20180101'  
GROUP BY TreeID;  
)
```

Show 4 ▾ entries

	TreeID	NumRipe	AvgTaste
1	1	2	3.5
2	2	7	3.71428571428571
3	3	3	3.33333333333333
4	4	3	3

Search:

Previous 1 2 3 4 5 Next 42 / 52

Showing 1 to 4 of 18 entries

It's up to you

```
tasty_bananas <- banana %>%
  dplyr::filter(DatePicked > '20180101') %>%
  dplyr::group_by(TreeID) %>%
  dplyr::summarise(NumRipe = n(), AvgTaste = mean(TasteRank, na.rm=T))
```

We can always inspect the SQL code created by `dplyr`.

```
tasty_bananas %>% show_query()

## <SQL>
## SELECT `TreeID`, COUNT(*) AS `NumRipe`, AVG(`TasteRank`) AS `AvgTaste` 
## FROM `Ape_Banana`
## WHERE (`DatePicked` > '20180101')
## GROUP BY `TreeID`
```

Advanced: programmatically edit queries

```
for (this_taste in c(3,4,5)) {  
  res <- DBI::dbGetQuery(con, stringr::str_interp(  
    "SELECT *  
    FROM Ape_Banana  
    WHERE TasterRank = ${this_taste}  
    "))  
  cat("\nResults for taste = ", this_taste, "\n")  
  print(nrow(res))  
}  
  
## Results for taste = 3  
## [1] 7  
## Results for taste = 4  
## [1] 10  
## Results for taste = 5  
## [1] 18
```

Advanced no.2: batch process results

Send the query without executing it.

```
rs <- DBI::dbSendQuery(con, "SELECT * FROM Ape_Banana")
```

Process results, 20 at a time

```
while (!DBI::dbHasCompleted(rs)) {
  twenty_bananas <- DBI::dbFetch(rs, n = 20)
  # << insert processing on twenty_bananas here >>
  print(nrow(twenty_bananas))
}

## [1] 20
## [1] 20
## [1] 10
```

Saving results

Collect the results

```
tasty_bananas <- dplyr::collect(tasty_bananas)
```

Save as a CSV ( **readr** is better at it).

```
library(readr)

# Choose the processed data directory
location <- here("data", "processed", "tasty_bananas.csv")

# Write CSV
readr::write_csv(tasty_bananas, location)
```

Saving results

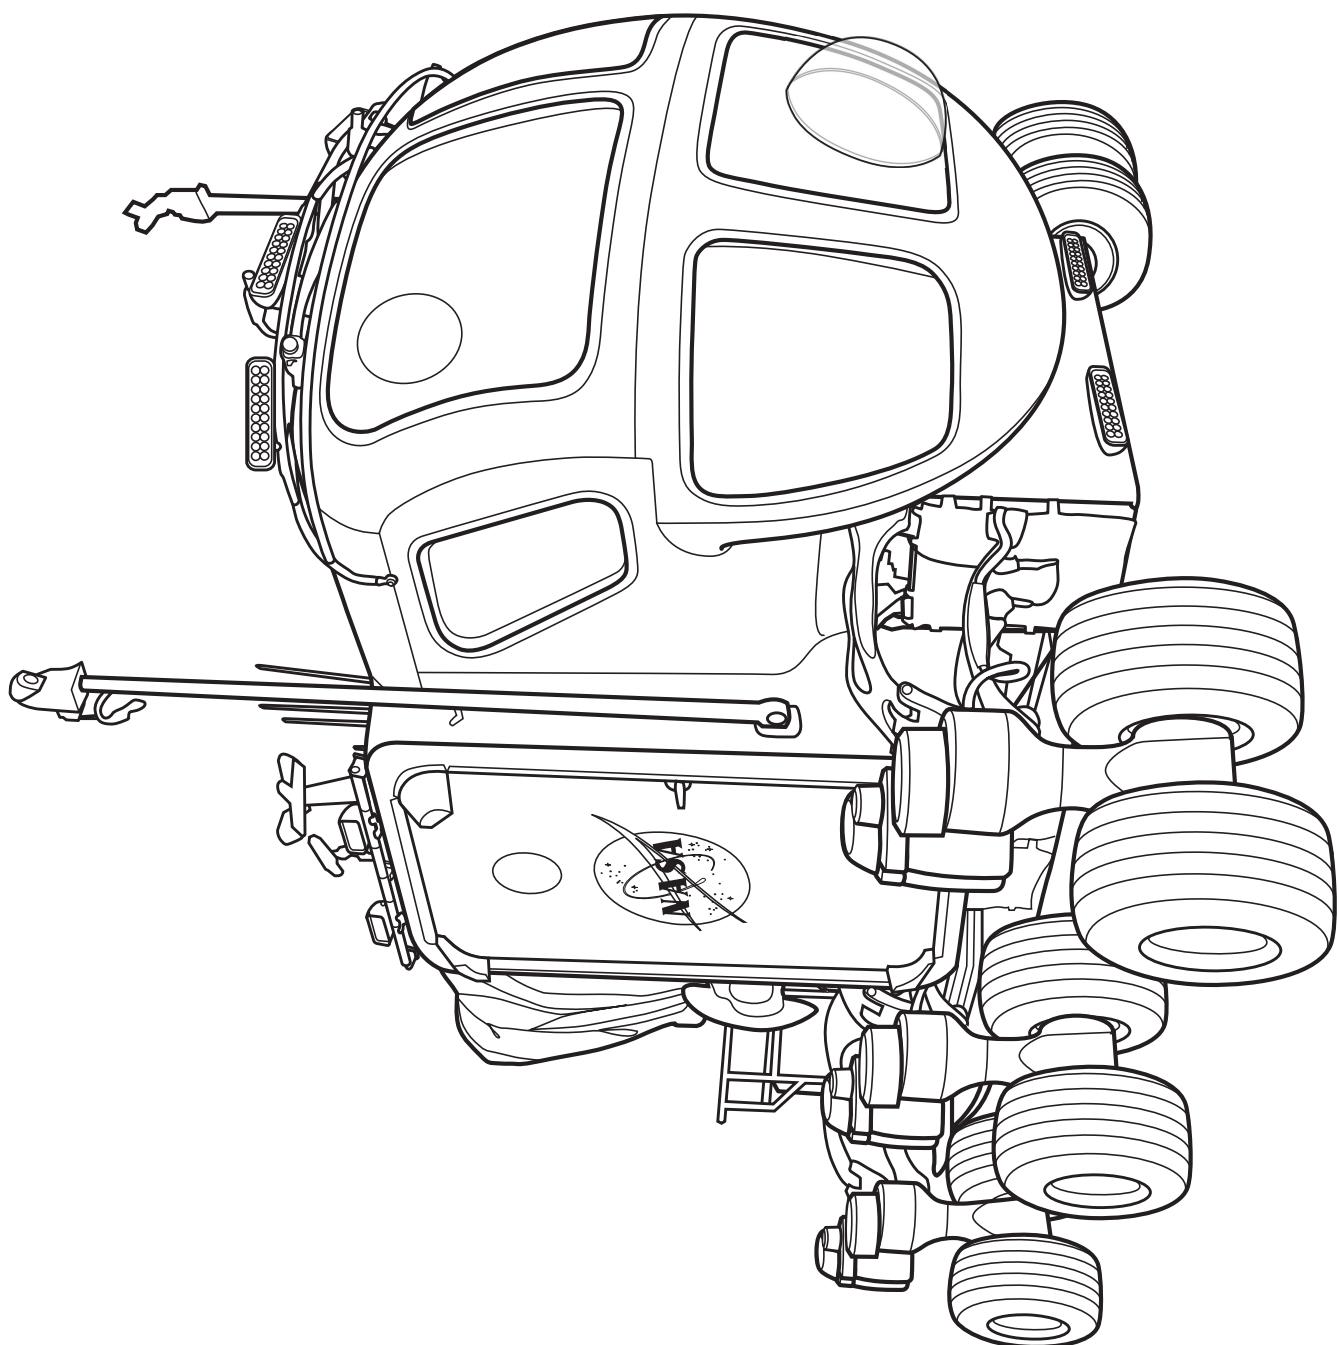
Or save to SQLite.

```
# Choose the processed data directory
location <- here("data", "processed", "Sandpit_results.sqlite")  
  
# Connect or create database
res_con <- DBI::dbConnect(RSQLite::SQLite(), location)  
  
# Save the table
DBI::dbWriteTable(res_con, "tasty_bananas", tasty_bananas)
```

Don't forget to disconnect

When you're done.

```
DBI::dbDisconnect(con)
DBI::dbDisconnect(res_con)
```



Live guide and demonstrations

Live guide and demonstrations

- Using GitHub
- Downloading **umbrella** repository
- Using **DB Browser** with SQLite
- Live demo of connecting and exploring

Also, see the **R code folder** on the course repo.

